



# Spatial prediction with frozen visual slots

Qwen3.5-4B visual tower | matched 20-epoch decoders | 4,096 training images


(a) Pseudo reference




(b) Human reference



(c) 4x super-resolution



(d) Matched 128 px crop



— Right lung — Left lung — Heart (pseudo only)

Inset: absolute HR error

